

Tasks:

1. Create a Kubernetes cluster on GCP. If possible, share a script / code which can be used to create the cluster.
➔ Created using Google Kubernetes Engine through Google Cloud UI
2. Install nginx ingress controller on the cluster. For now, we consider that the user will add public IP of ingress LoadBalancer to their /etc/hosts file for all hostnames to be used. So do not worry about DNS resolution.
➔ Ingress controller needs a specific namespace, service account, cluster role bindings, configmaps etc. One can create all the Kubernetes objects mentioned using the yaml file from the official ingress repo.
 1. `$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/master/deploy/static/mandatory.yaml`
 2. `$ kubectl apply -f ingress-nginx_service.yaml`

Default response of nginx after installation, listening on public IP address: 35.222.129.198



Also updated /etc/hosts file on local machine

```
archit@archit-dell: ~  
archit@archit-dell:~$ cat /etc/hosts  
127.0.0.1      localhost  
#127.0.1.1    archit-dell  
35.222.129.198 guestbook.mstakx.io  
35.222.129.198 staging-guestbook.mstakx.io
```

3. On this cluster, create namespaces called staging and production.
 1. `$ kubectl create namespace staging`
namespace/staging created
 2. `$ kubectl create namespace production`
namespace/production created
3. Install [guest-book](#) application on both namespaces.
 1. Installing guest-book application in staging namespace
`$ kubectl create -f redis-master-deployment.yaml -n staging`
`$ kubectl create -f redis-master-service.yaml -n staging`

```
$ kubectl create -f redis-slave-deployment.yaml -n staging
$ kubectl create -f redis-slave-service.yaml -n staging
$ kubectl create -f frontend-deployment.yaml -n staging
$ kubectl create -f frontend-service_nginx.yaml -n staging
```

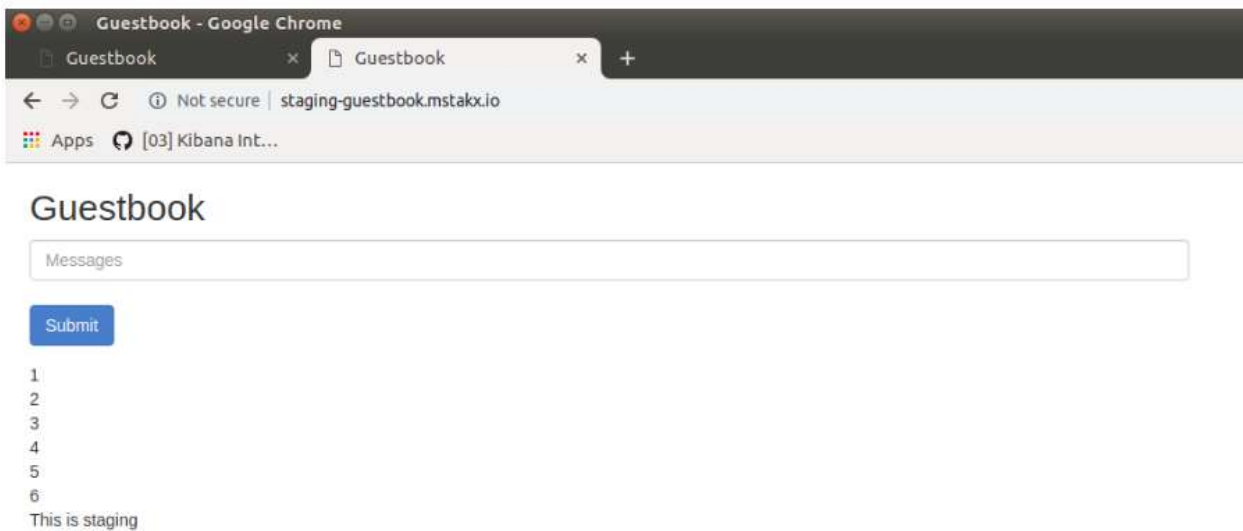
2. Installing guest-book application in staging namespace

```
$ kubectl create -f redis-master-deployment.yaml -n production
$ kubectl create -f redis-master-service.yaml -n production
$ kubectl create -f redis-slave-deployment.yaml -n production
$ kubectl create -f redis-slave-service.yaml -n production
$ kubectl create -f frontend-deployment.yaml -n production
$ kubectl create -f frontend-service_nginx.yaml -n production
```

- **Note:** frontend-deployment.yaml has little change then the one mentioned on github, added auto scale part and reduced CPU requirement, so that can auto scale pod with lower traffic.

3. Expose staging application on hostname staging-guestbook.mstakx.io

1. \$ kubectl create -f ingress_staging_guestbook_routing.yaml -n staging



2. Expose production application on hostname `guestbook.mstakx.io`
 1. `$ kubectl create -f ingress_production_guestbook_routing.yaml -n production`



Messages

1
2
3
4
5
This is production

State of Kubernetes' pod, service and namespace after executing all above steps

```

architmehta06@cloudshell:~ (storied-reserve-243808)$ kubectl get svc --all-namespaces
NAMESPACE      NAME                TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
default        kubernetes          ClusterIP    10.59.240.1      <none>           443/TCP          21h
ingress-nginx  ingress-nginx       LoadBalancer 10.59.248.178    35.222.129.198   80:31556/TCP,443:32632/TCP 20h
kube-system    default-http-backend NodePort     10.59.247.143    <none>           80:31972/TCP      21h
kube-system    heapster            ClusterIP    10.59.244.1      <none>           80/TCP           21h
kube-system    kube-dns            ClusterIP    10.59.240.10     <none>           53/UDP,53/TCP     21h
kube-system    metrics-server      ClusterIP    10.59.242.165    <none>           443/TCP           21h
production     frontend            ClusterIP    10.59.254.21     <none>           8088/TCP          11h
production     redis-master        ClusterIP    10.59.255.185    <none>           6379/TCP          11h
production     redis-slave         ClusterIP    10.59.244.10     <none>           6379/TCP          11h
staging        frontend            ClusterIP    10.59.244.62     <none>           8088/TCP          11h
staging        redis-master        ClusterIP    10.59.251.218    <none>           6379/TCP          11h
staging        redis-slave         ClusterIP    10.59.247.146    <none>           6379/TCP          11h
architmehta06@cloudshell:~ (storied-reserve-243808)$ kubectl get ingress --all-namespaces
NAMESPACE      NAME                HOSTS                ADDRESS          PORTS          AGE
production     production-ingress  guestbook.mstakx.io  35.222.129.198  80            89m
staging        staging-ingress     staging-guestbook.mstakx.io 35.222.129.198  80            90m
architmehta06@cloudshell:~ (storied-reserve-243808)$

```

```
architmehta06@cloudshell:~ (storied-reserve-243808)$ kubectl get ingress --all-namespaces
NAMESPACE   NAME                 HOSTS                 ADDRESS          PORTS   AGE
production  production-ingress   guestbook.mstakx.io  35.222.129.198  80      89m
staging      staging-ingress       staging-guestbook.mstakx.io  35.222.129.198  80      90m
architmehta06@cloudshell:~ (storied-reserve-243808)$
```

```

architmehta06@cloudshell:~ (storied-reserve-243808)$ kubectl get pods --all-namespaces
NAMESPACE          NAME                                     READY   STATUS    RESTARTS   AGE
ingress-nginx      nginx-ingress-controller-76c86d76c4-hvjt6  1/1     Running   0           21h
kube-system        event-exporter-v0.2.4-5f7d5d7dd4-sjk4r    2/2     Running   0           21h
kube-system        fluentd-gcp-scaler-7b895cbc89-s7tc7       1/1     Running   0           21h
kube-system        fluentd-gcp-v3.2.0-6czt6                  2/2     Running   0           21h
kube-system        fluentd-gcp-v3.2.0-ppgnr                  2/2     Running   0           21h
kube-system        fluentd-gcp-v3.2.0-wxlck                   2/2     Running   0           21h
kube-system        heapster-v1.6.0-beta.1-64486f697-lw6p5    3/3     Running   0           21h
kube-system        kube-dns-autoscaler-76fcd5f658-czmkz       1/1     Running   0           21h
kube-system        kube-dns-b46cc9485-21l29                  4/4     Running   0           21h
kube-system        kube-dns-b46cc9485-mcvqp                   4/4     Running   0           21h
kube-system        kube-proxy-gke-test-default-pool-552a9cf4-4vfh  1/1     Running   0           21h
kube-system        kube-proxy-gke-test-default-pool-552a9cf4-skss  1/1     Running   0           21h
kube-system        kube-proxy-gke-test-default-pool-552a9cf4-zhmm  1/1     Running   0           21h
kube-system        l7-default-backend-6f8697844f-5ghrh        1/1     Running   0           21h
kube-system        metrics-server-v0.3.1-5b4d6d8d98-ghgrb     2/2     Running   0           21h
kube-system        prometheus-to-sd-9r7cn                     1/1     Running   0           21h
kube-system        prometheus-to-sd-ct2pc                     1/1     Running   0           21h
kube-system        prometheus-to-sd-rpvz2                     1/1     Running   0           21h
production        frontend-84bb688cb6-mtg6n                  1/1     Running   0           11h
production        redis-master-57fc67768d-9kccg              1/1     Running   0           11h
production        redis-slave-7556d5fd6-w6qdx                1/1     Running   0           11h
staging            frontend-84bb688cb6-vz6jr                  1/1     Running   0           11h
staging            redis-master-57fc67768d-g22tn              1/1     Running   0           11h
staging            redis-slave-7556d5fd6-4lk4c                1/1     Running   0           11h

```

- Implement a pod autoscaler on both namespaces which will scale frontend pod replicas up and down based on CPU utilization of pods.
 - ➔ Added "HorizontalPodAutoscaler" kind with minReplicas: 1 and maxReplicas: 3 with targetCPUUtilizationPercentage: 3 , this component added to "frontend-deployment.yaml"
- Write a script which will demonstrate how the pods are scaling up and down by increasing/decreasing load on existing pods.
 - ➔ Wrote sample shell script to send http request to front end in infinite loop, and started watch on "kubectl get pod" command in other terminal.

Script: load_generator.sh (added to github)

```
#!/bin/bash
cnt=0;
while true; do
    register=$(curl -v http://guestbook.mstakx.io/ )
    cnt=`expr $cnt + 1`
done
echo $cnt;
```

Below snapshot shows that when script was running, after some time it created another pod

ScaleUp

```
architmehta06@cloudshell:~ (storied-reserve-243808)$ kubectl top pod frontend-77c48c5988-6kztj -n production
NAME                                CPU (cores)   MEMORY (bytes)
frontend-77c48c5988-6kztj           1m            10Mi
architmehta06@cloudshell:~ (storied-reserve-243808)$ kubectl get pods -n production -w
NAME                                READY   STATUS    RESTARTS   AGE
frontend-77c48c5988-6kztj           1/1     Running   0           68m
redis-master-57fc67768d-9kccg       1/1     Running   0           16h
redis-slave-7556d5fd6-w6qdx         1/1     Running   0           16h
frontend-77c48c5988-grxgh           0/1     Pending   0           0s
frontend-77c48c5988-grxgh           0/1     Pending   0           0s
frontend-77c48c5988-grxgh           0/1     ContainerCreating   0           0s
frontend-77c48c5988-grxgh           1/1     Running   0           2s
^Carchitmehta06@cloudshell:~ (storied-reserve-243808)$ kubectl get pods -n production
NAME                                READY   STATUS    RESTARTS   AGE
frontend-77c48c5988-6kztj           1/1     Running   0           70m
frontend-77c48c5988-grxgh           1/1     Running   0           47s
redis-master-57fc67768d-9kccg       1/1     Running   0           16h
redis-slave-7556d5fd6-w6qdx         1/1     Running   0           16h
```

ScaleDown: After that, I stopped script and it automatically deleted the 2nd pod after sometime

```
architmehta06@cloudshell:~ (storied-reserve-243808)$ kubectl get pods -n production
NAME                                READY   STATUS    RESTARTS   AGE
frontend-77c48c5988-6kztj           1/1     Running   0           83m
frontend-77c48c5988-grxgh           1/1     Running   0           14m
redis-master-57fc67768d-9kccg       1/1     Running   0           17h
redis-slave-7556d5fd6-w6qdx         1/1     Running   0           17h
architmehta06@cloudshell:~ (storied-reserve-243808)$ kubectl get pods -n production
NAME                                READY   STATUS    RESTARTS   AGE
frontend-77c48c5988-6kztj           1/1     Running   0           93m
redis-master-57fc67768d-9kccg       1/1     Running   0           17h
redis-slave-7556d5fd6-w6qdx         1/1     Running   0           17h
```

- Write a wrapper script which does all the steps above. Mention any pre-requisites in the README.md at the root of your repo.

The evaluator will proceed by going over the steps mentioned in the README. So try to make this as automated as possible.

 - ➔ Remaining : But we can have a python script which does

1. Installation steps for installing nginx and guestbook application for list of namespaces
2. Updating etc hostFile based on ExternIP for LoadBalancer
3. Sending curl to staging and production to verify that it is up and running and validating response

In the context of above test, please explain the following:

- What was the node size chosen for the Kubernetes nodes? And why?
➔ Number of nodes chosen: 3 (1 Master 2 Worker)
Just went with default 3 node cluster. That is enough for the demo, for production we need to have backup of master node as and worker node depends on the compute and application requirement.
Machine Type: n1-standard-1 (1 vCPU, 3.75 GB memory) which makes total 3 vCPU and memory 11.25 GB
While choosing, goal was just to implement this particular assignment. I looked at basic guestbook application for which “cpu: 50m” and “memory: 100Mi” with replicacount 1 and not being cpu or memory intensive this quota is sufficient.
- What method was chosen to install the demo application and ingress controller on the cluster, justify the method used
➔ Got all the yaml files for guestbook Application, nginxcontroller and adding ingress object, created separate yaml files for deployment and service for better understanding. Used “kubectl create” to apply all yaml files.
This way one can clearly understand the steps and process about how complete application is getting set up and working. Alternative way is to prepare helm chart and with single command we can deploy whole set up.
- What would be your chosen solution to monitor the application on the cluster and why?
➔ ELK/EFK: Application is used to solve business problem and the logging events are also generated according to business logic which are highly dimensional. EFK uses indexing mechanism to store everything and provides option to search. Having to search/monitor on n number of key this works efficiently. It is also industry proven and heavily used with large scale.
- What additional components/plugins would you install on the cluster to manage it better?
➔ prometheus : To measure cluster health and other metrics (ex: container_cpu_usage_total)
Grafana: for viewing matrices
Both are opensource part of CNCF and scalable.